

# CS4495/6495

## Introduction to Computer Vision

---

8B-L1 *Classification: Generative models*




# *Supervised* classification


Given a collection of labeled examples, come up with a function that will predict the labels of new examples.

Training examples

“four”



“nine”



Novel input

?

# Supervised classification

How good is the function we come up with to do the classification? (What does “good” mean?)

Depends on:

- What mistakes does it make
- Cost associated with the mistakes

# Supervised classification

Since we know the desired labels of training data, we want to *minimize the expected misclassification*

# Supervised classification

## Two general strategies

- Use the training data to build representative probability model; separately model class-conditional densities and priors (*Generative*)
- Directly construct a good decision boundary, model the posterior (*Discriminative*)

# Supervised classification: Generative

Given labeled training examples, predict labels for new examples

- Notation:  $(4 \rightarrow 9)$  - object is a '4' but you call it a '9'
- We'll assume the cost of  $(X \rightarrow X)$  is zero.

# Supervised classification: Generative

Consider the two-class (binary) decision problem:

- $L(4 \rightarrow 9)$ : Loss of classifying a 4 as a 9
- $L(9 \rightarrow 4)$ : Loss of classifying a 9 as a 4

# Supervised classification: Generative


**Risk** of a classifier strategy **S** is expected loss:

$$R(S) = \Pr(4 \rightarrow 9 \mid \text{using } S) L(4 \rightarrow 9) \\ + \Pr(9 \rightarrow 4 \mid \text{using } S) L(9 \rightarrow 4)$$

We want to choose a classifier so as to minimize this total risk



# Supervised classification: minimal risk



At best decision boundary, either choice of label yields same expected loss.

Feature value  $x$


If we choose class “four” at boundary, expected loss is:

$$= P(\text{class is 9} | \mathbf{x}) L(9 \rightarrow 4) + P(\text{class is 4} | \mathbf{x}) L(4 \rightarrow 4)$$

If we choose class “nine” at boundary, expected loss is:

$$= P(\text{class is 4} | \mathbf{x}) L(4 \rightarrow 9)$$

# Supervised classification: minimal risk



Feature value  $x$

At best decision boundary, either choice of label yields same expected loss.


So, best decision boundary is at point  $x$  where:

$$P(\text{class is 9} | x) L(9 \rightarrow 4) = P(\text{class is 4} | x) L(4 \rightarrow 9)$$

To classify a new point, choose class with lowest expected loss; i.e., choose

“four” if:  $P(4 | x) L(4 \rightarrow 9) > P(9 | x) L(9 \rightarrow 4)$

# Supervised classification: minimal risk



At best decision boundary, either choice of label yields same expected loss.

So, best decision boundary is at point  $x$  where:


$$P(\text{class is } 9|x) L(9 \rightarrow 4) = P(\text{class is } 4|x) L(4 \rightarrow 9)$$

*How to evaluate these probabilities?*

# Example: learning skin colors




$P(x | \text{skin})$



Feature  $x = \text{Hue}$

$P(x | \text{not skin})$




# Example: learning skin colors




Now we get a new image,  
and want to label each pixel  
as skin or non-skin.


$P(x | \text{skin})$



$P(x | \text{not skin})$



# Bayes rule



$$\overbrace{P(\text{skin} \mid x)}^{\text{posterior}} = \frac{\overbrace{P(x \mid \text{skin})}^{\text{likelihood}} \overbrace{P(\text{skin})}^{\text{prior}}}{P(x)}$$

$$P(\text{skin} \mid x) \propto P(x \mid \text{skin}) \boxed{P(\text{skin})}$$

*Where does the prior come from?*

# Bayes rule in (ab)use

Likelihood ratio test (assuming cost of errors is the same):

If  $P(\text{skin}|\mathbf{x}) > P(\sim\text{skin}|\mathbf{x})$  classify  $\mathbf{x}$  as skin

... SO ....

If  $P(\mathbf{x}|\text{skin})P(\text{skin}) > P(\mathbf{x}|\sim\text{skin})(P(\sim\text{skin}))$   
classify  $\mathbf{x}$  as skin (Bayes rule)

*(if the costs are different just re-weight)*

# Bayes rule in (ab)use

... but I don't really know prior  $P(\textit{skin})$ ...

... but I can assume it some constant  $\Omega$  ...

... so with some training data I can **estimate**  $\Omega$  ...

.... and with the same training data I can **measure** the **likelihood densities** of **both**  $P(x|\textit{skin})$  and  $P(x|\sim\textit{skin})$  ...

So.... I can more or less come up with a rule...



# Example: classifying skin pixels

Now for every pixel in a new image, we can estimate probability that it is generated by skin:

If  $p(\text{skin}|x) > \theta$  classify as skin; otherwise not



Brighter pixels are  
higher probability  
of being skin

# Example: classifying skin pixels




**Figure 6:** A video image and its flesh probability image




**Figure 7:** Orientation of the flesh probability distribution marked on the source video image

# Example: classifying skin pixels



**Figure 13:** CAMSHIFT-based face tracker used to “fly” over a 3D graphic’s model of Hawaii



**Figure 12:** CAMSHIFT-based face tracker used to play Quake 2 hands free by inserting control variables into the mouse queue


# More general generative models

For a given measurement  $\mathbf{x}$  and set of classes  $c_i$   
choose  $c^*$  by:

$$c^* = \arg \max_c p(c | \mathbf{x}) = \arg \max_c p(c) p(\mathbf{x} | c)$$

# Continuous generative models

- If  $\mathbf{x}$  is continuous, need *likelihood* density model of  $p(\mathbf{x}|c)$
- Typically parametric – Gaussian or mixture of Gaussians



# Continuous generative models

- Why not just some histogram or some KNN (Parzen window) method?
  - You might...
  - But you would need lots and lots of data everywhere you might get a point
  - The whole point of modeling with a parameterized model is not to need lots of data.

# Summary of generative models:

- + Firm probabilistic grounding
- + Allows inclusion of prior knowledge
- + *Parametric modeling of likelihood permits using small number of examples*
- + *New classes do not perturb previous models*
- + Others:
  - Can take advantage of unlabelled data
  - Can be used to generate samples

## Summary of generative models:

- And just where did you get those priors?
- Why are you modeling those obviously non-C points?
- The example hard cases aren't special
- If you have lots of data, doesn't help



## Next...

- A really cool way of building a generative model for face recognition (not detection)
- And then *discriminative* models...