

CS4495/6495

Introduction to Computer Vision

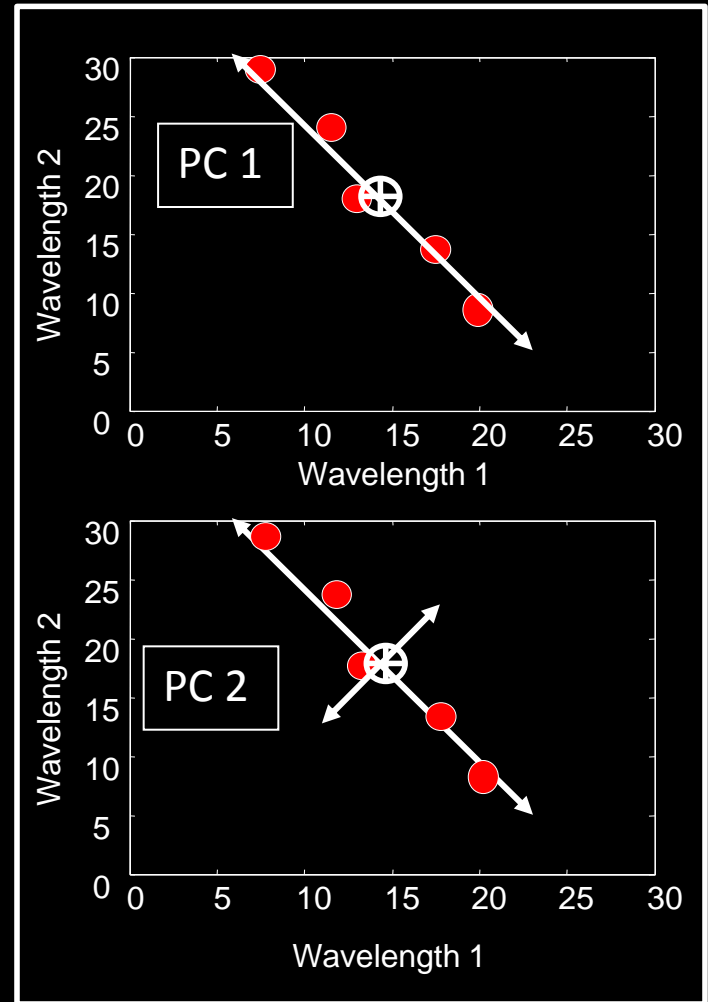
8B-L2 *Principle Component Analysis
(and its use in Computer Vision)*



Figure 2. Seven of the eigenfaces calculated from the input images of Figure 1.


Principal Components

- *Principal components* are all about the directions in a feature space along which points have the greatest variance.



Principal Components

- First PC is the direction of maximum variance. Technically (and mathematically) it's from the origin, but we actually mean the mean.
- Subsequent PCs are orthogonal to previous PCs and describe maximum residual variance




2D example: Fitting a line

$$E(a, b, d) = \sum_i (ax_i + by_i - d)^2$$

$$\frac{\partial E}{\partial d} = 0 \rightarrow -2 \sum_i (ax_i + by_i - d) = 0$$

$$d = a\bar{x} + b\bar{y}$$




2D example: Fitting a line

Substitute $d = a\bar{x} + b\bar{y}$:

$$E = \sum_i [a(x_i - \bar{x}) + b(y_i - \bar{y})]^2 = \|\mathbf{B}n\|^2$$

where $\mathbf{B} = \begin{pmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ x_2 - \bar{x} & y_2 - \bar{y} \\ \dots & \dots \\ x_n - \bar{x} & y_n - \bar{y} \end{pmatrix}$


so minimize $\|\mathbf{B}n\|^2$
subject to $\|n\|=1$ gives
axis of least inertia



Sound familiar???

Another algebraic interpretation

Minimizing sum of squares of distances to the line is the same as *maximizing the sum of squares of the projections on that line*.



Algebraic interpretation

Trick: How is the sum of squares of projection lengths expressed in algebraic terms?

Line

p	p	p	...	p
t	t	t	...	t
1	2	3	...	m

Point 1
Point 2
Point 3
:
Point m

L
i
n
e

 x^T B^T B x

$$= \sum_i (\vec{x}^T P_i)^2$$

Algebraic interpretation

Our goal: $\max(\mathbf{x}^T \mathbf{B}^T \mathbf{B} \mathbf{x})$, subject to $\mathbf{x}^T \mathbf{x} = 1$

Line

p	p	p	...	p
t	t	t	...	t
1	2	3	...	m

Point 1
Point 2
Point 3
:
Point m

L
i
n
e

 \mathbf{x}^T \mathbf{B}^T \mathbf{B} \mathbf{x}

$$= \sum_i (\vec{x}^T P_i)^2$$

Algebraic interpretation

$$\max(\mathbf{x}^T \mathbf{B}^T \mathbf{B} \mathbf{x}), \text{ subject to } \mathbf{x}^T \mathbf{x} = 1$$

$$\text{maximize } E = \mathbf{x}^T \mathbf{M} \mathbf{x} \text{ subject to } \mathbf{x}^T \mathbf{x} = 1 \quad (\mathbf{M} = \mathbf{B}^T \mathbf{B})$$

$$E' = \mathbf{x}^T \mathbf{M} \mathbf{x} + \lambda(1 - \mathbf{x}^T \mathbf{x})$$

$$\frac{\partial E'}{\partial \mathbf{x}} = 2\mathbf{M} \mathbf{x} + 2\lambda \mathbf{x}$$

$$\frac{\partial E'}{\partial \mathbf{x}} = 0 \rightarrow \mathbf{M} \mathbf{x} = \lambda \mathbf{x} \quad (\mathbf{x} \text{ is an } \textit{eigenvector} \text{ of } \mathbf{B}^T \mathbf{B})$$

Yet another algebraic interpretation

$$\mathbf{B}^T \mathbf{B} = \begin{pmatrix} \sum_{i=1}^n x_i^2 & \sum_{i=1}^n x_i y_i \\ \sum_{i=1}^n x_i y_i & \sum_{i=1}^n y_i^2 \end{pmatrix} \text{ if about } \textit{origin}$$

So the principal components are the orthogonal directions of the covariance matrix of a set points.

Yet one more algebraic interpretation

$$\mathbf{B}^T \mathbf{B} = \sum \mathbf{x} \mathbf{x}^T \quad \text{if about } \textit{origin}$$

$$\mathbf{B}^T \mathbf{B} = \sum (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \quad \textit{otherwise} - \textit{outer product}$$


So the principal components are the orthogonal directions of the *covariance* matrix of a set points.

Eigenvectors


How many eigenvectors are there?

- For Real Symmetric Matrices of size $N \times N$
 - Except in degenerate cases when eigenvalues repeat, there are N distinct eigenvectors

PCA: Eigenvalues



Dimensionality Reduction



We can represent the **orange** points with *only* their v_1 coordinates

Higher Dimensions

- Suppose each data point is N-dimensional

$$\text{var}(\mathbf{v}) = \sum_{\mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2$$

$$= \mathbf{v}^T \mathbf{A} \mathbf{v} \text{ where } \mathbf{A} = \sum (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T$$

Outer product



- The *ev* with largest eigenvalue λ captures the most variation among training vectors \mathbf{x}
 - eigenvector with smallest eigenvalue has least variation

The space of all face images

When viewed as *vectors* of pixel values, face images are extremely high-dimensional

- 100x100 image = 10,000 dimensions



The space of all face images


However, relatively few
10,000-dimensional vectors
correspond to valid face
images

- We want to effectively
model the subspace of face
images



The space of all face images

We want to construct a *low-dimensional linear subspace* that best explains the variation in the set of face images



Principal Component Analysis

- Given: M data points x_1, \dots, x_M in R^d where d is big
- We want some directions in R^d that capture most of the variation of the x_i . The coefficients would be:

$$u(x_i) = u^T(x_i - \mu)$$

(μ : mean of data points)

- What unit vector u in R^d captures the most variance of the data?

Principal Component Analysis

- Direction that maximizes the variance of the projected data:

$$\text{var}(\mathbf{u}) = \frac{1}{M} \sum_{i=1}^M \underbrace{\mathbf{u}^T (\mathbf{x}_i - \mu)(\mathbf{u}^T (\mathbf{x}_i - \mu))^T}_{\text{Projection of data point}}$$

Projection of data point

$$= \mathbf{u}^T \left[\underbrace{\frac{1}{M} \sum_{i=1}^N (\mathbf{x}_i - \mu)(\mathbf{x}_i - \mu)^T}_{\text{Covariance matrix of data}} \right] \mathbf{u}$$

Covariance matrix of data

$$= \mathbf{u}^T \Sigma \mathbf{u}$$

Principal Component Analysis

- Direction that maximizes the variance of the projected data:

$$\text{var}(\mathbf{u}) = \mathbf{u}^T \Sigma \mathbf{u}$$

Since \mathbf{u} is a unit vector that can be expressed in terms of some linear sum of the eigenvectors of Σ , then direction of \mathbf{u} that maximizes the variance is the eigenvector associated with the largest eigenvalue of Σ

Principal component analysis

- The *direction that captures the maximum covariance* of the data is the *eigenvector* corresponding to the *largest eigenvalue* of the *data covariance* matrix
- Furthermore, the top k orthogonal directions that capture the most variance of the data are the k eigenvectors corresponding to the k largest eigenvalues
- But first, we'll need the PCA d>>>n trick...

The dimensionality trick

Let Φ_i be the (very big vector length d) that is face image I with the mean image subtracted.

Define $C = \frac{1}{M} \sum \Phi_i \Phi_i^T = AA^T$

where $A = [\Phi_1 \Phi_2 \dots \Phi_M]$ is the matrix of faces, and is $d \times M$.

Note: C is a huge $d \times d$ matrix (remember d is the length of the vector of the image).

The dimensionality trick

So $C = AA^T$ is a huge matrix.

But consider $A^T A$. It is only $M \times M$. Finding those eigenvalues is easy.

Suppose \mathbf{v}_i is an eigenvector $A^T A$:

$$A^T A \mathbf{v}_i = \lambda \mathbf{v}_i$$

Premultiply by A :

$$AA^T A \mathbf{v}_i = \lambda A \mathbf{v}_i$$

So: $A \mathbf{v}_i$ are the eigenvectors of $C = AA^T$

How many eigenvectors are there?

- If had $M > d$ then there would be d . But $M \ll d$.
- So intuition would say there are M of them.
- But wait: if 2 points in 3D, how many eigenvectors? 3 points?
- Subtracting out the mean yields $M-1$.

Eigenfaces: Key idea *(Turk and Pentland, 1991)*

- Assume that most face images lie on a low-dimensional subspace determined by the first k ($k \ll d$) directions of maximum variance
- Use PCA to determine the vectors or “eigenfaces” u_1, u_2, \dots, u_k that span that subspace
- Represent all face images in the dataset as linear combinations of eigenfaces. Find the coefficients by dot product.

Eigenfaces example


Training
images

$$\mathbf{x}_1, \dots, \mathbf{x}_M$$



Top eigenvectors: u_1, \dots, u_k

Mean: μ



Eigenfaces example

Principal component (eigenvector) u_k



Eigenfaces example

Principal component (eigenvector) u_k



$\mu + 3\sigma_k u_k$



$\mu - 3\sigma_k u_k$



Eigenfaces example

- Face \mathbf{x} in “face space” coordinates (dot products) :



$$\mathbf{x} \rightarrow [\mathbf{u}_1^T (\mathbf{x} - \boldsymbol{\mu}), \dots, \mathbf{u}_k^T (\mathbf{x} - \boldsymbol{\mu})] \\ = [w_1, \dots, w_k]$$

This vector is the representation of the face.

Eigenfaces example

- Reconstruction:



The diagram illustrates the reconstruction of a face image using Eigenfaces. It shows a target face image (labeled \hat{x}) being reconstructed as the sum of a mean face image (labeled μ) and a weighted sum of eigenface images (labeled $u_1, u_2, u_3, u_4, \dots$).

$$\hat{x} = \mu + w_1 u_1 + w_2 u_2 + w_3 u_3 + w_4 u_4 + \dots$$

Eigenfaces example

- But what about recognition?

Recognition with eigenfaces

Given novel image \mathbf{x} :

- Project onto subspace:

$$[w_1, \dots, w_k] = [u_1^T(\mathbf{x} - \mu), \dots, u_k^T(\mathbf{x} - \mu)]$$

- Optional: check reconstruction error $\mathbf{x} - \hat{\mathbf{x}}$ to determine whether image is really a face
- Classify as closest training face in k-dimensional subspace
- This is why it's a *generative* model.

And old cast of characters...



Limitations – PCA of global structure

- Global appearance method: not robust to misalignment, background variation



Limitations – PCA in general

- The direction of maximum variance is not always good for classification

